

# Introduction to R

Andy Grogan-Kaylor

June 29, 2019

## Contents

1	Background	2
2	Introduction	2
3	Base R and Libraries	3
4	Working Directory	3
5	Writing R Code or Script	4
6	Graphical User Interface	4
7	Get Your Data	4
7.1	R format (*.RData)	5
7.2	Comma Separated Values (*.csv)	5
7.3	Statistical Packages and Excel	5
7.4	foreign	5
7.5	readxl	5
7.6	haven	6
8	Save Your Data in R Format	6
9	Save and Document Your Work	6
10	Process Your Data	6
10.1	Data Subsets	6
10.2	Numeric and Factor Variables	7
10.3	Missing Values	7
10.4	Renaming Variables	8
10.5	Sorting Data	8
10.6	Creating New Variables	8
10.7	Scales or Measures	8
11	Descriptive Statistics	9
12	Bivariate Statistics	9
12.1	Crosstabulation	9
12.2	Correlation	9
12.3	t Test	10
12.4	ANOVA	10

13	Multivariate Statistics	10
14	Graphing	10
15	Comments, Questions and Corrections	11

## 1 Background

This guide is mainly written for academics, community based researchers, and advocates, who are interested in using R to analyze and visualize data.

R has a number of advantages for individuals working in academic settings, agencies, and community settings. First of all because R is open source, R is free, and does not have a high cost like proprietary statistical software or data visualization software.

Second, using R means that one has access to a worldwide community of people who are constantly developing new R packages, and new materials for learning R.

That being said, R can have a number of drawbacks. Documentation and help files can sometimes be difficult to understand. R's syntax, and the "R way of doing things" can present a formidable barrier.

My hope in this document is to provide an introduction to R that bypasses some of these difficulties by providing straightforward instruction focused on the likely needs of social researchers, community based researchers, and advocates. I want to help these groups of people to use R in an effective way.

I believe that it is possible to teach R in an accessible way, and that a little bit of R can take you a long way.

## 2 Introduction

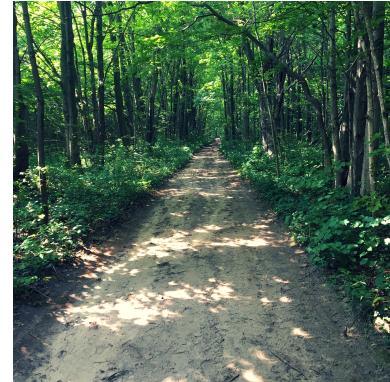
This work is licensed under a [Creative Commons Attribution-ShareAlike 4.0 International License](#)

This document is a brief introduction to R<sup>1</sup>. Commands that you actually type into R are represented in `courier font`. `mydata` is the name of your data set. `x` and `y` and `z` refer to variables in your data. More documentation on any command is usually available via `help(command)` or `??command`. The R interface makes it extremely easy to do rapid interactive data analysis. Hit "**Up-Arrow**" to recall the most recent command, which you can then quickly edit and resubmit. Remember also that one often submits a command or set of commands from a script window.

The general idea of many R commands is:

```
command(data = mydata, ...variables..., options)
```

or



## Introduction to R

A Practically Focused Guide

Andrew Grogan-Kaylor

agrogan@umich.edu | www.umich.edu/~agrogan

Figure 1: Cover Image

<sup>1</sup> This document is inspired by my long-standing "Two Page Stata" document: ([PDF](#)) ([HTML](#)).

```
command(mydata$xvar, options)
```

The \$ sign is a kind of “connector”. mydata\$x means: “The variable x in the dataset called mydata”.

Sometimes, it is not necessary to use any options since some authors of R have done a good job of thinking about the defaults. R can make use of long pathnames to files like:

```
C:/Users/user1/Desktop/mydata.sav
```

Note that R uses forward slashes / instead of backslashes for directories. R uses ~ to refer to the user's (usually your) home directory.

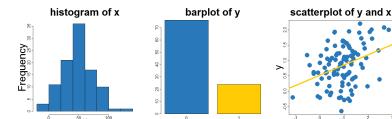


Figure 2: Graphical Possibilities in Base R

### 3 Base R and Libraries

Most of this guide makes use of what is most often called **Base R**, the R that you get when you install the R software, and RStudio, on your computer.

For many social researchers, the data structure of primary interest is the *data frame*, and thus that is my focus here. In the interests of parsimony I do not go into a great deal of detail on R's other data structures.

A great deal can be accomplished with **Base R**. However, as you grow in your use of R, you will likely frequently need to make use of libraries, which are invoked by the library( . . . ) command. Before using a library you need to install it. Below is an example of installing the *ggplot2* advanced graphics library.

You would need to install the library only once. Installation can also be accomplished from the “Packages” tab in RStudio.

```
install.packages("ggplot2")
```

Then start the library when you are using R by typing...

```
library(ggplot2)
```

I should mention here the new additions to the R language of the new libraries which make up the **tidyverse**. Learning the **tidyverse** requires an additional investment in learning, however the **tidyverse** makes many improvements to the R language and functionality.

### 4 Working Directory

R uses the concept of a *working directory* to know where to find files, and where to save files.

It is often helpful to simply set your working directory to a particular location and by default, files will be accessed from, and saved to, that directory e.g.:

```
getwd() # "get", or find out, your working directory
setwd("C:/Users/user1/Desktop/") # set your working directory
```

Note that R uses a forward slash / to specify directory paths. R does not understand the use of a backward slash \ to specify directories. R uses ~ to refer to the user's (usually your) home directory.

## 5 Writing R Code or Script

R is a command or syntax based program, and many advanced functions are only available via syntax.

R Commands are stored in a *script* or *code* file that usually ends in .R, e.g. myRscript.R. The command file is distinct from your actual data, stored in an .RData file, e.g. mydata.RData.

Base R can sometimes be cryptic.

However, a little bit of Base R can go a long way, and you can get a great learning return for a little bit of investment in learning Base R.

## 6 Graphical User Interface

A good **Graphical User Interface** (GUI) can make some of the base functionality of R available without the use of syntax. **RCommander** is the best GUI, and can be installed from the command line by typing:

```
install.packages("Rcmdr", dependencies=TRUE)
```

**RCommander** can make some tasks easier, but the syntax that it produces can sometimes be non-intuitive. Often it is easiest (and more in the interests of replicable research) just to learn how to write the R code that accomplishes a particular task. Further, your learning may go quicker if you bypass **RCommander** altogether and simply learn how to write R code.

**RStudio** is an **Integrated Development Environment (IDE)** that can be run simultaneously with **RCommander** and provides an easier working environment for R Software. If all the software is installed, Start **RStudio** to start R, then type library(Rcmdr) to start **RCommander**.

## 7 Get Your Data

Remember that R uses a forward slash / to specify directory paths. R does not understand the use of a backward slash \ to specify directories.

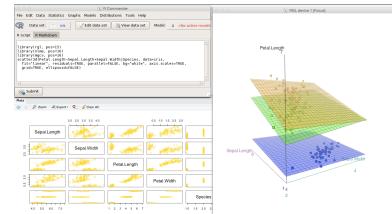


Figure 3: screenshot of RCommander

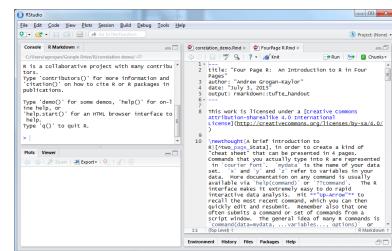


Figure 4: screenshot of RStudio

## 7.1 R format (\*.RData)

R most easily makes use of data in R format. Data can be loaded with the `load()` command.

```
load("the/path/to/myRfile.RData") # specific directory path and file

load("myRfile.RData") # no path indicated; file needs to be in working directory
```

Note—as we discuss in a little more detail below—that a single data file can contain multiple data frames.

For example, a data file called **projectdata.RData** could contain:

- A data frame on *clients*, called **clients**.
- A data frame on *providers*, called **providers**.
- A data frame on *facilities*, called **facilities**.

The name of the *RData file* can be **very different** from the name of the *data frames* that it contains.

## 7.2 Comma Separated Values (\*.csv)

R can also read *comma separated values (csv)*.

```
library(readr) # to read csv

mydata <- read_csv("myCSVfile.csv")
```

## 7.3 Statistical Packages and Excel

R can easily import well-formatted data from other packages<sup>2</sup> like SPSS, Stata, or Excel<sup>2</sup>.

## 7.4 foreign

```
library(foreign) # library for importing from stats software

mydata <- read.spss("mySPSSfile.sav") # SPSS

mydata <- read.dta("myStatafile.dta") # Stata
```

<sup>2</sup> These instructions assume you have `setwd()` appropriately, or alternatively are specifying a full pathname and filename.

## 7.5 readxl

```
library(readxl) # library for importing Excel files

mydata <- read_excel("mySpreadsheet.xls") # Excel
```

## 7.6 haven

haven is a new library for reading in data from statistical packages that may work better.

```
library(haven) # new library for importing from stats software

dataset <- read_sav("mySPSSfile.sav") # SPSS

dataset <- read_stata("myStatafile.dta") # Stata
```

## 8 Save Your Data in R Format

Once you have your data in R, it will likely make sense to save it in \*.RData format for future use.

```
save(mydata, file = "mydata.RData")
```

Note—as we alluded to earlier—that multiple data frames can be saved into a single data file.<sup>3</sup>

```
save(clients, # a first data frame, about clients
     providers, # a second data frame, about providers
     facilities, # a third data frame, about facilities
     file = "projectdata.RData")
```

<sup>3</sup> Some would call this a *feature* of R, while others would simply say that this is another *confusing aspect* of R.

## 9 Save and Document Your Work

Use the Script Editor to save R commands that you want to use again, or to modify for the next project, as well as to create an “audit trail” of your work so that your workflow is documented and replicable. R commands are saved in a .R file, e.g. **myscript.R**.

## 10 Process Your Data

### 10.1 Data Subsets

Working with a *subset* of your data (i.e. fewer variables rather than many many variables) is often helpful. The subset function can be especially helpful.

```
mydata_subset <- subset(mydata, # name of data
                        age > 18, # condition(s)
                        select = c(id, sex, income)) # variables
```

You can then run functions like `summary()` on a *subset* of your data.

```
summary(mydata_subset)
```

You can also save this subset for future use.

```
save(mydata_subset, file = "mydata_subset.RData")
```

## 10.2 Numeric and Factor Variables

R recognizes two basic kinds of variables: *continuous variables* (which R calls “*numeric*” variables) which are often scales like income, mental health, or neighborhood quality; and *categorical variables* (which R calls “*factor variables*”) like race, gender or religion.

R seems to make a stronger distinction between these two types of variables than some other statistical software.<sup>4</sup> It can thus sometimes be useful to change variables from one type to another:

```
mydata$x <- as.numeric(mydata$y)
```

```
mydata$x <- as.factor(mydata$y) # shorter syntax
```

```
# longer, more complete syntax
```

```
mydata$w <- factor(mydata$z, # original numeric variable
                    levels = c(0, 1, 2), # levels of numeric variable
                    labels = c("Group A", # labels
                              "Group B",
                              "Group C"),
                    ordered = TRUE) # often useful to order the levels
```

## 10.3 Missing Values

Data with missing values, often represented as negative numbers (e.g. -99, -9, -8) need to be recoded so that the missing values are represented as a missing value character (“NA”) that R knows to exclude from calculations.

```
mydata$x[mydata$x == -9] <- NA # Example 1
```

```
mydata$x[mydata$x == -8] <- NA # Example 2
```

Sometimes you want to drop rows of data that contain missing values. This can be accomplished with `na.omit()`.

```
mydata2 <- na.omit(mydata)
```

`na.omit()` removes a row of data where *any* value is missing, so sometimes you want to work with a subset of your data before applying `na.omit()`.

```
mydata_subset <- subset(mydata, # name of data
                        age > 18, # condition(s)
```

<sup>4</sup> In many cases, this is very helpful in that R recognizes that the type of variable calls for a certain kind of statistic or graph, or vice versa. In other cases, this may be the source of an *error message*.

```

select = c(id, sex, income) # variables

mydata_subset2 <- na.omit(mydata_subset)

```

## 10.4 Renaming Variables

It is often convenient to rename your data so that the variables have more intuitively understandable names e.g.

```

mydata$age <- mydata$var123

mydata$gender <- mydata$var456

```

## 10.5 Sorting Data

It is sometimes useful to sort your data. `sort(mydata$x)` will sort mydata by the values of x.

## 10.6 Creating New Variables

You can easily create new variables in R. For example, a change score between a measure collected at two time-points, like a pre-test, and a post-test, would be:

```
mydata$change_x <- mydata$xTime2 - mydata$xTime1
```

## 10.7 Scales or Measures

Similarly, you can sum the items of a scale into a scale as follows:

```
mydata$myscale <- mydata$x1 + mydata$x2 + mydata$x3
```

You can test the alpha reliability of this scale with the following syntax:

```
myscale_data <- subset(mydata, select = c(x1, x2, x3))
```

The syntax above create a dataframe of only the scale items.

Then,

```
library(psych)
```

```
alpha(myscale_data)
```

## 11 Descriptive Statistics

`summary(mydata$x)` gives you basic descriptive statistics for a variable, such as the mean (average). Especially useful for continuous variables. Use `summary(mydata)` to summarize every variable in your data.

`skim(mydata)` from `library(skimr)` or `describe(mydata)` from `library(psych)` will often give you a nicer summary of your variables that is closer to what you want for an academic paper or agency report.

You may only want to look at descriptive statistics for a subset of your data.

Creating a `subset` and then running descriptive statistics on that subset may be helpful.

`table(mydata$x)` gives you a frequency distribution for your variable. Especially useful for *factor variables*. `prop.table(table(mydata$x))` will give you a table of proportions. Calling up `library(descr)` and then using `freq(mydata$x)` will give you a more nicely formatted frequency distribution.

## 12 Bivariate Statistics

### 12.1 Crosstabulation

Tabulating two categorical variables (*factor variables*) together gives you a cross-tabulation of those variables, e.g:

```
table(mydata$x, mydata$y) # simple table of counts

prop.table(table(mydata$x, mydata$y)) # table of cell proportions

prop.table(table(mydata$x, mydata$y)
           margin = 1) # row margins: row proportions

prop.table(table(mydata$x, mydata$y)
           margin = 2) # column margins: column proportions
```

then

```
chisq.test(table(mydata$x, mydata$y))
```

will give you a chi-square test of the relationship of x and y.

### 12.2 Correlation

```
cor(mydata[, c("x", "y")], use="complete.obs")
```

will give you the correlation of continuous variables x and y<sup>5</sup>.

<sup>5</sup> Here is an example where R turns a simple issue into a difficult one, and the syntax is frankly less than elegant, and not intuitive. I don't have this syntax memorized. I use `library(Rcmdr)` if I need to test, or create a script for, a correlation.

```
cor.test(mydata$x, mydata$y,
         alternative="two.sided",
         method="pearson")
```

will test the statistical significance of this correlation.

### 12.3 t Test

```
numSummary(mydata$x, groups=mydata$z)
```

gives you a summary of continuous variable x by *factor variable* z.

```
t.test(mydata$x~mydata$z)
```

runs a t-test of continuous variable x over *factor variable* z.

### 12.4 ANOVA

```
aov(x ~ z, data=mydata)
```

runs the corresponding ANOVA of continuous variable x across *factor variable* z.

## 13 Multivariate Statistics

```
mymodel <- lm(y ~ x + z, data=mydata)
```

runs a regression (linear model) of y on x and z.

Type `summary(mymodel)` to display the results.

## 14 Graphing

```
hist(mydata$x)
```

will give you a nice display of one continuous variable.

```
hist(mydata$x, main="...", xlab="...")
```

gives a nicer looking graph.

```
barplot(table(mydata$x))
```

gives similar results when x is a *factor variable*.

```
plot(mydata$y, mydata$x)
```

gives you a twoway scatterplot of your data

A more nicely labelled graph can be obtained with:

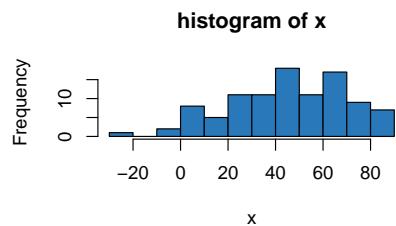


Figure 5: histogram of x

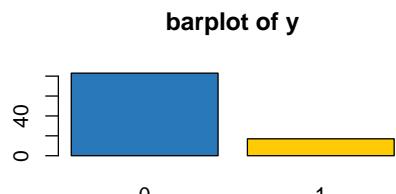


Figure 6: barplot of y

```
plot(y, x,
  main= "...",
  xlab= "...",
  ylab= "...")
```

`abline(lm(mydata$y~mydata$x))` will add a linear fit line to a scatter-plot that you have already constructed.

`abline(lm(mydata$y~mydata$x), col="gold", lwd=5)` will be a nicer looking fit line.

## 15 Comments, Questions and Corrections

Comments, questions and corrections most welcome and may be sent to:

Andrew Grogan-Kaylor @ <http://www.umich.edu/~agrogan> & @ [agrogan@umich.edu](mailto:agrogan@umich.edu).

Last updated: June 29 2019 at 13:14

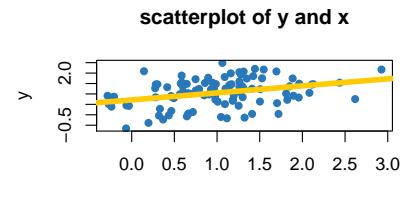


Figure 7: scatterplot of y against x